Descriptive and Exploratory Analysis

Jason G. Fleischer, Ph.D.
Asst. Teaching Professor
Department of Cognitive Science, UC San Diego
jfleischer@ucsd.edu


@jasongfleischer

https://jgfleischer.com

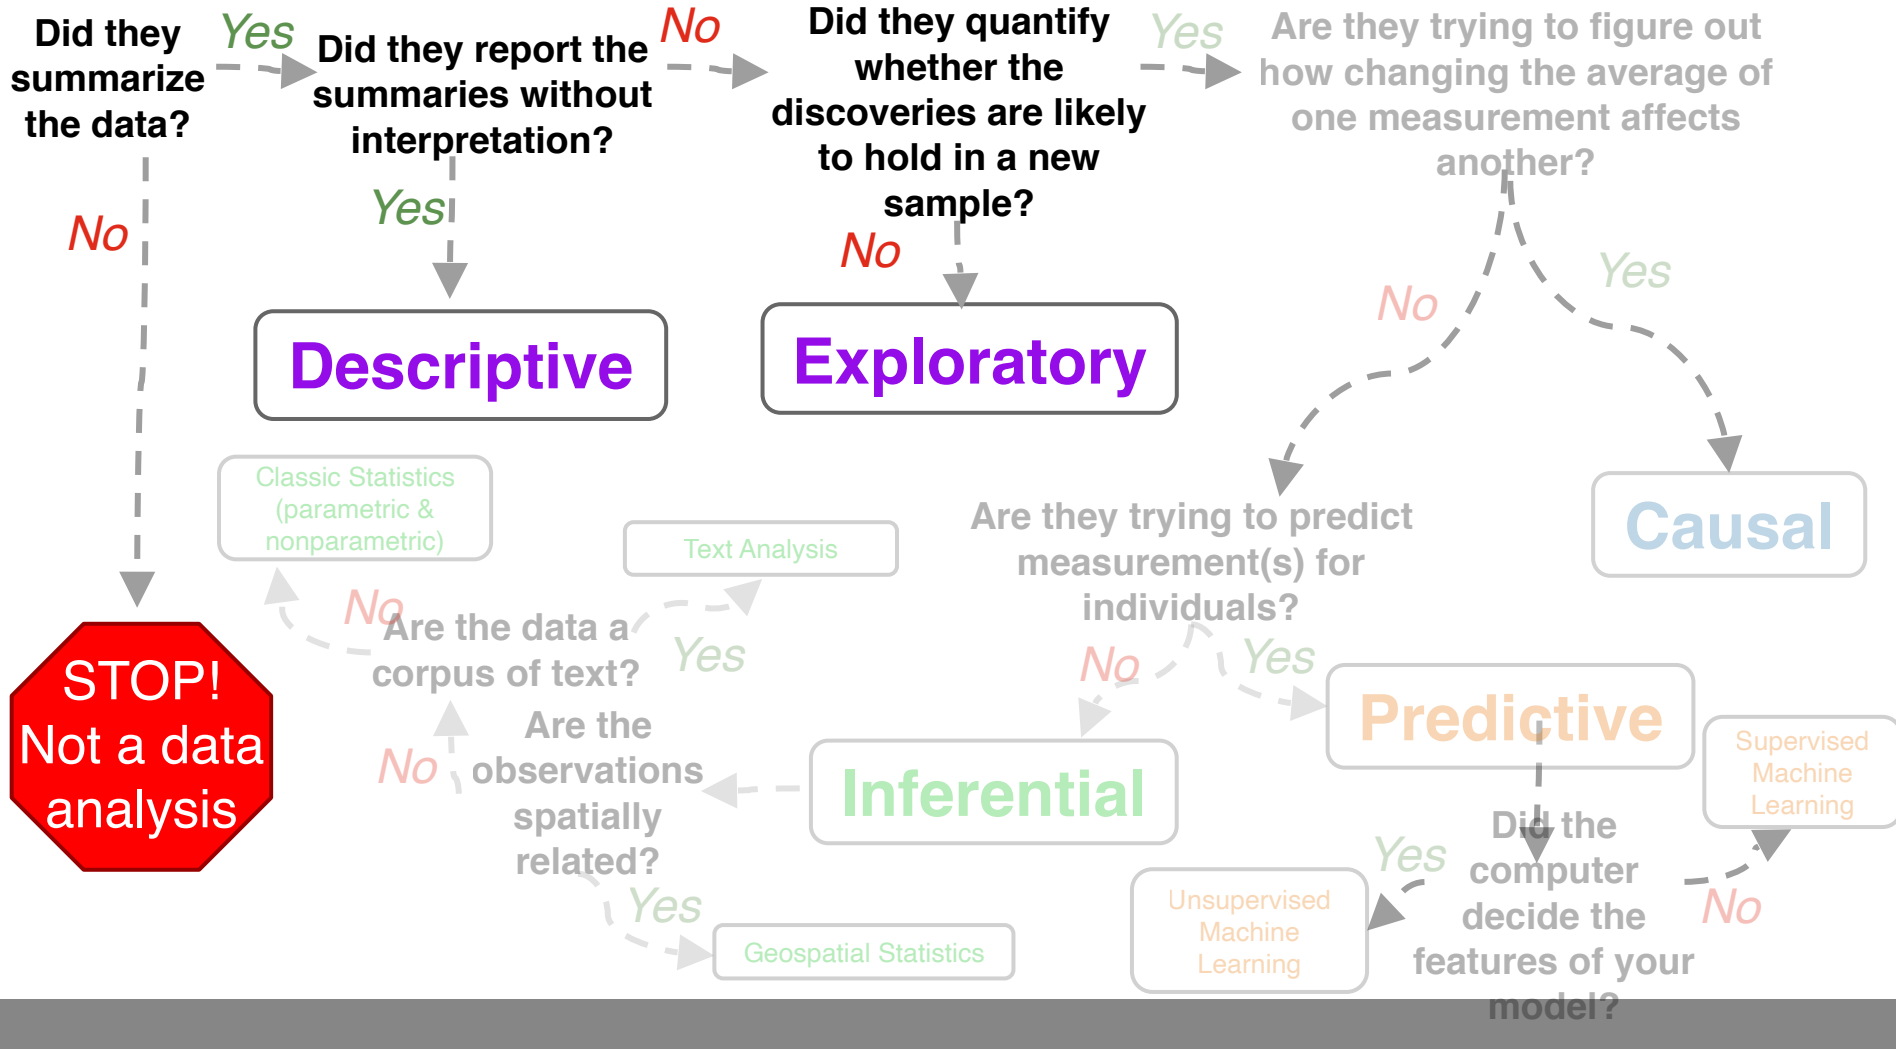
Did they summarize the data?
Yes
Did they report the summaries without interpretation?
No
Did they quantify whether the discoveries are likely to hold in a new sample?
Yes
Are they trying to figure out how changing the average of one measurement affects another?
No
Yes
No
Yes
No
Yes
Descriptive
Exploratory
Causal
Classic Statistics (parametric & nonparametric)
Text Analysis
Are they trying to predict measurement(s) for individuals?
No
Yes
No
Yes
STOP! Not a data analysis
Are the data a corpus of text?
Predictive
Supervised Machine Learning
Are the observations spatially related?
Inferential
Yes
Did the computer decide the features of your model?
No
Yes
Unsupervised Machine Learning
Geospatial Statistics

Descriptive: The goal of descriptive analysis is to understand the components of a data set, describe what they are, and explain that description to others who might want to understand the data.

- Problem: Understanding whether users are nice or mean on Youtube
- Data science question: Are the words that people use in their comments more frequently positive words (great, awesome, nice, useful) or negative words (bad, stupid, lame, awful)?
- Type of analysis: Descriptive analysis

To answer this you would calculate statistics about YouTube comments

Statistics

*“the science that deals with the **collection, classification, analysis, and interpretation of numerical facts or data**”*

Source: dictionary.com

statistic

“*A quantity computed from a sample*”

Source: dictionary.com

Populations & Samples

We want to learn something about this..

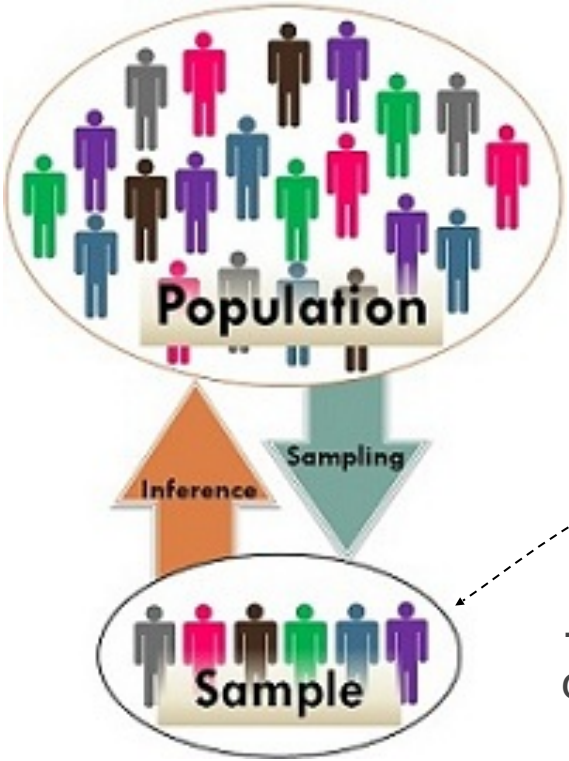

Our <u>population</u>: *all* YouTube comments

Our <u>sample</u>: 100,000 comments

....but we can only *actually* collect data from this

Image source: keydifferences.com

statistic

"*A quantity computed from a sample*"

For our YouTube analysis, we could take a random sample of comments from YouTube and calculate the following statistic: *the number of positive and the number of negative words in each review*.

Source: dictionary.com

Best sampling practices:

- Always think about what your population is
- Collect data from a sample that is representative of your population
- If you have no choice but to work with a dataset that is not collected randomly and is biased, be careful not to generalize your results to the entire population

You'd want to be sure you sample randomly across *all* YouTube comments, making sure not to get more comments from one genre over another, or one location over another, etc.

Examples of bad sampling:

- Surveying subscribers of a gun-related magazine for research on Americans' attitudes toward owning guns
- Randomly sampling Facebook users for what TV shows people like

To understand *all* YouTube comments, you wouldn't just want to sample from one YouTube channel, or videos in a single language.

It's *always* worth spending time at the beginning of a project to determine whether or not the data you have are garbage. Be certain they are actually able to help you answer the question you're interested in.

GIGO : Garbage In. Garbage Out.

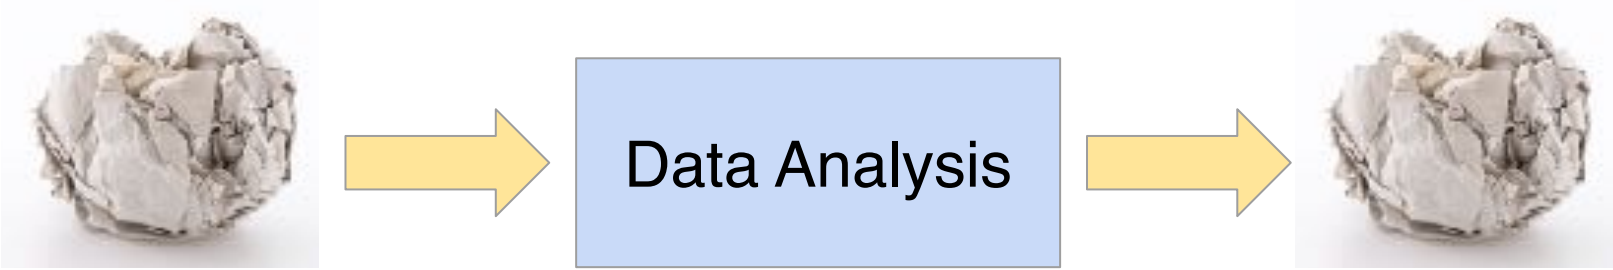

For the survey data I collected from you all, which of the following best describes the population I could generalize findings back to.

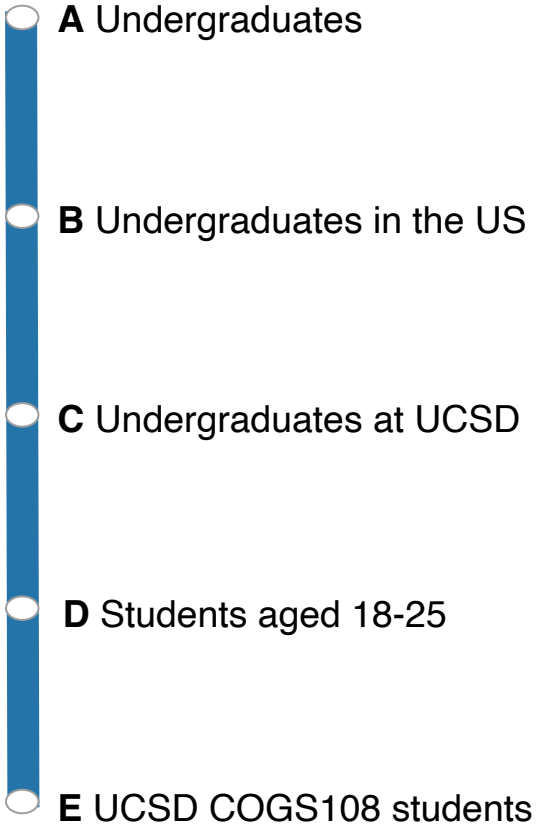

Descriptive

Descriptive Analysis

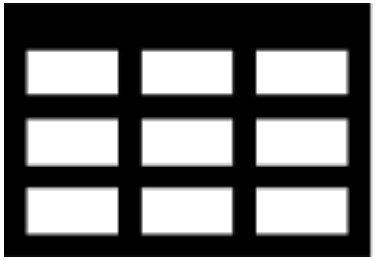

Size

Missingness

Shape

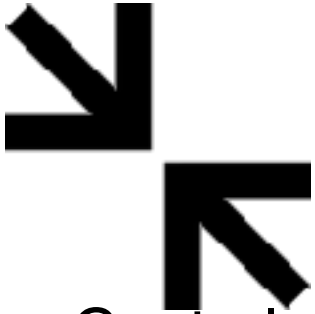

Central Tendency

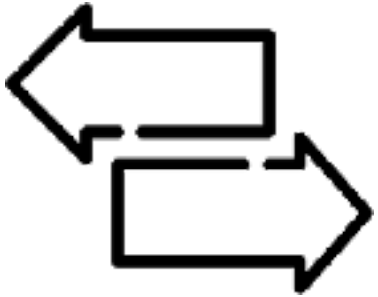

Variability

Descriptive

Size

How many observations (rows) and variables (columns) you have is an important first step. You should always be aware of the size of your dataset

Descriptive

Missingness

It's critical to know how many observations have missing data for variables of interest in your data. Knowing *why* their missing is also important.

Descriptive

Shape

It’s critical to know the distribution of the variables in your dataset. Certain statistical approaches can only be used with certain distributions.

Descriptive

Central Tendency

Knowing the mean, median, and/or mode can help you get an idea of what a typical value is for your variable(s) of interest

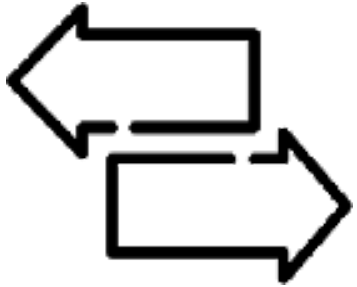

Variability

The central tendency tells you part of the story. The variability in the values in your observation helps fill in the rest.

Which of the following is NOT something accomplished by a descriptive analysis?

A Describes typical values in your dataset

B Determines the size of your dataset

C Establishes causal relationships between variables

D Identifies missing data

E Determines how variable values in your dataset are

Descriptive Statistics & Summary

"*We must suppress some of the truth to communicate the truth... In short, the techniques of descriptive statistics are designed to match the salient features of the data set to human cognitive abilities.*"

-I.J. Good (1983)

Descriptive

Descriptive Analyses are often included as “Table 1” in academic publications

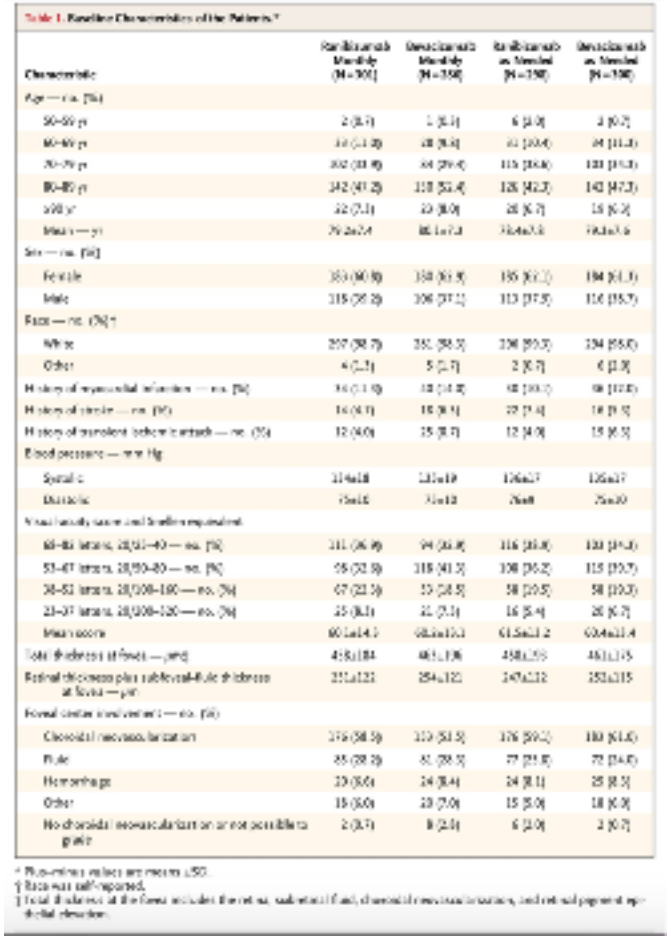

Characteristic				
Age — no. (%)				
50–59 yr	2 (0.7)	1 (0.3)	6 (2.0)	2 (0.7)
60–69 yr	33 (11.0)	28 (9.3)	41 (13.6)	34 (11.3)
70–79 yr	102 (33.9)	94 (31.3)	115 (38.3)	121 (40.3)
80–89 yr	142 (47.2)	153 (51.0)	126 (42.0)	142 (47.3)
≥90 yr	22 (7.3)	23 (8.0)	12 (4.0)	11 (3.7)
Mean — yr	79.2±7.4	80.1±7.3	78.4±6.5	78.8±6.6
Sex — no. (%)				
Female	183 (60.8)	168 (56.0)	155 (51.7)	184 (61.3)
Male	118 (39.2)	132 (44.0)	145 (48.3)	116 (38.7)
Race — no. (%)†				
White	297 (98.7)	292 (97.3)	298 (99.3)	294 (98.0)
Other	4 (1.3)	8 (2.7)	2 (0.7)	6 (2.0)
History of myocardial infarction — no. (%)	34 (11.3)	43 (14.3)	42 (14.0)	36 (12.0)
History of stroke — no. (%)	16 (5.3)	18 (6.0)	22 (7.3)	18 (6.0)
History of transient ischemic attack — no. (%)	12 (4.0)	15 (5.0)	12 (4.0)	15 (5.0)
Blood pressure — mm Hg				
Systolic	134±18	135±19	136±17	135±17
Diastolic	73±10	73±10	76±9	75±10
Visual acuity, snellen or ETDRS equivalent				
68–83 letters; 20/32–40 — no. (%)	111 (36.9)	99 (33.0)	116 (38.7)	120 (39.3)
53–67 letters; 20/50–80 — no. (%)	95 (32.5)	118 (41.3)	108 (36.2)	115 (38.7)
38–52 letters; 20/100–160 — no. (%)	67 (22.3)	53 (18.5)	58 (19.5)	58 (19.3)
23–37 letters; 20/200–320 — no. (%)	25 (8.3)	21 (7.3)	16 (5.4)	20 (6.7)
Mean score	60.1±14.3	58.4±15.1	61.5±13.2	60.4±13.4
Total thickness at fovea — μm‡	458±184	481±196	458±159	461±175
Retinal thickness plus subfoveal choroidal thickness at fovea — μm	251±121	254±121	247±112	252±115
Foveal center involvement — no. (%)				
Choroidal neovascularization	176 (58.7)	153 (53.5)	176 (59.2)	180 (61.0)
Fluid	85 (28.3)	81 (28.3)	77 (25.9)	72 (24.4)
Hemorrhage	23 (7.7)	24 (8.4)	24 (8.1)	25 (8.5)
Other	15 (5.0)	20 (7.0)	15 (5.0)	18 (6.1)
No choroidal neovascularization or not accessible to grade	2 (0.7)	8 (2.8)	6 (2.0)	2 (0.7)

* Plus–minus values are means ±SD.
† Race was self-reported.
‡ Total thickness at the fovea includes the retina, subretinal fluid, choroidal neovascularization, and retinal pigment epithelial elevation.

https://www.nejm.org/doi/full/10.1056/nejmoa1102673

Descriptive

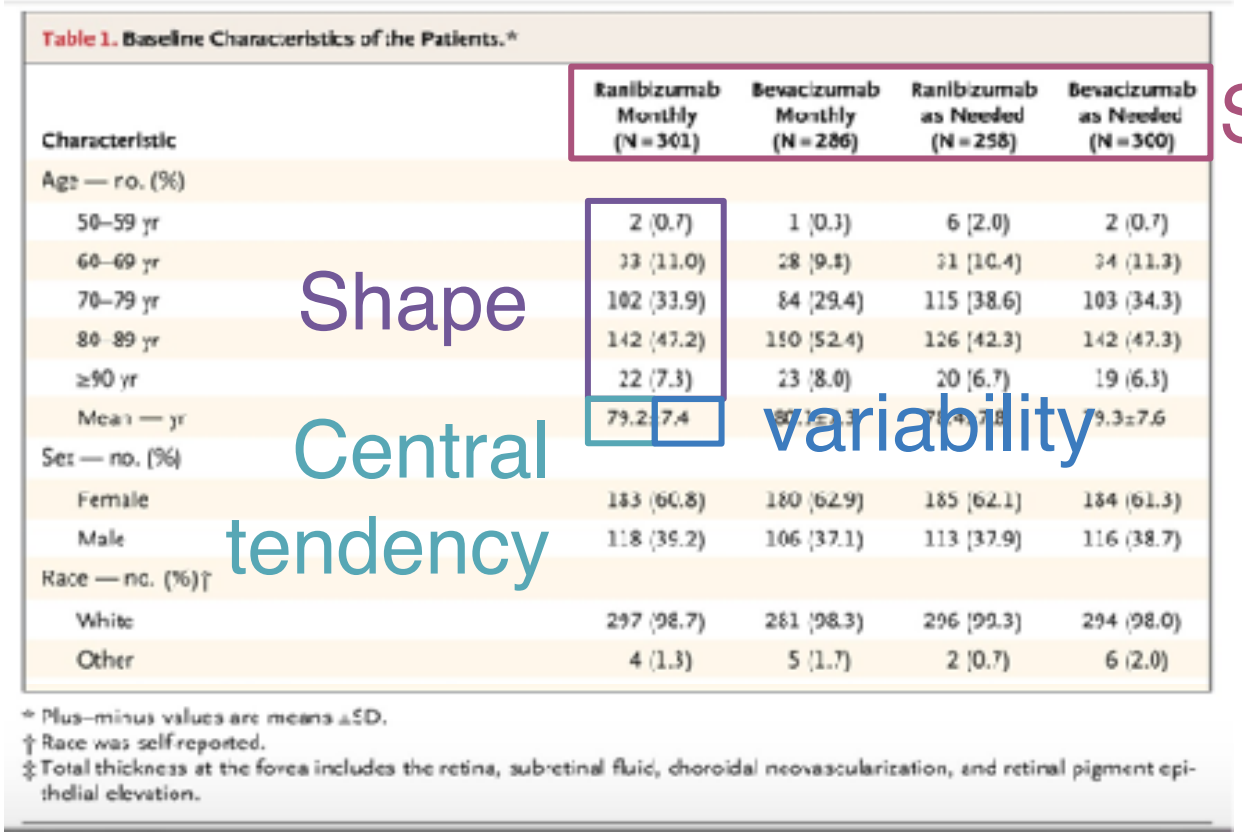

Table 1. Baseline Characteristics of the Patients.*

Characteristic	Ranibizumab Monthly (N=301)	Bevacizumab Monthly (N=286)	Ranibizumab as Needed (N=258)	Bevacizumab as Needed (N=300)
Age — no. (%)				
50–59 yr	2 (0.7)	1 (0.3)	6 (2.0)	2 (0.7)
60–69 yr	33 (11.0)	28 (9.8)	31 (10.4)	34 (11.3)
70–79 yr	102 (33.9)	84 (29.4)	115 (38.6)	103 (34.3)
80–89 yr	142 (47.2)	150 (52.4)	126 (42.3)	142 (47.3)
≥90 yr	22 (7.3)	23 (8.0)	20 (6.7)	19 (6.3)
Mean — yr	79.2±7.4	[illegible]	[illegible]	79.3±7.6
Sex — no. (%)				
Female	183 (60.8)	180 (62.9)	185 (62.1)	184 (61.3)
Male	118 (39.2)	106 (37.1)	113 (37.9)	116 (38.7)
Race — no. (%)†				
White	297 (98.7)	281 (98.3)	296 (99.3)	294 (98.0)
Other	4 (1.3)	5 (1.7)	2 (0.7)	6 (2.0)

* Plus–minus values are means ±SD.
† Race was self-reported.
‡ Total thickness at the fovea includes the retina, subretinal fluid, choroidal neovascularization, and retinal pigment epithelial elevation.

Size

Shape

Central tendency

variability

Zooming in on this we see variables stratified by Age, Sex, and Race

https://www.nejm.org/doi/full/10.1056/nejmoa1102673

Descriptive

Descriptive Statistics & Summary

Calculating descriptive statistics, understanding what they tell you about your data, and reporting them are critical steps in every analysis.

Exploratory: The goal is to find unknown relationships between the variables you have measured in your data set. Exploratory analysis is open ended and designed to verify expected or find unexpected relationships between measurements.

Exploratory

Exploratory Data Analysis (EDA) detective work answering the question: "*What can the data tell us?*"

Exploratory

Why EDA?

- Understand data properties
- Discover Patterns
- Generate & Frame Hypothesis
- Suggest modeling strategies
- Check assumptions (sanity checks)
- Communicate results (present the data)

…..and if you don’t, you’ll regret it

Exploratory

The dataset

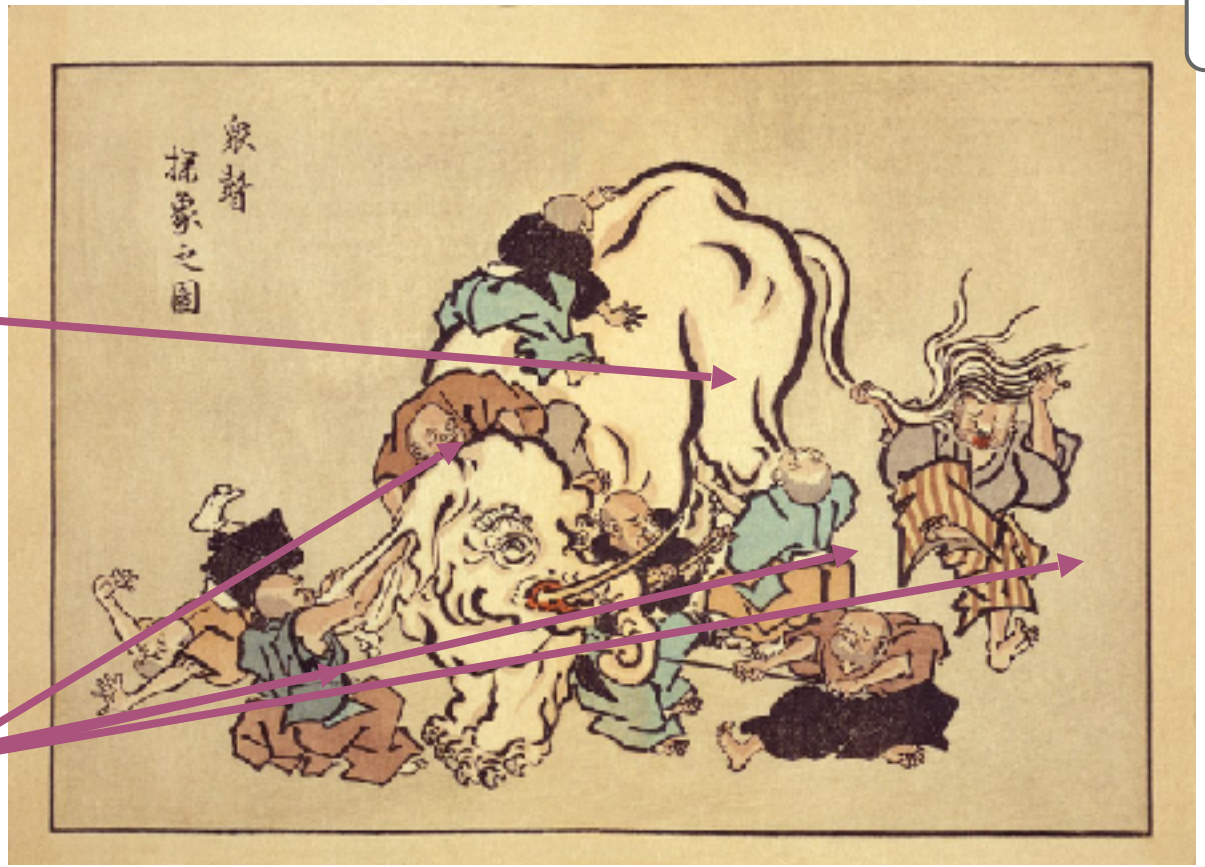

You

Exploratory

The general principles of exploratory analysis:

- Look for missing values
- Look for outlier values
- Calculate numerical summaries
- Generate plots to explore relationships
- Use tables to explore relationships
- If necessary, transform variables

EDA Approaches to “Get a Feel for the Data”

Understanding the relationship between variables in your dataset

Exploratory

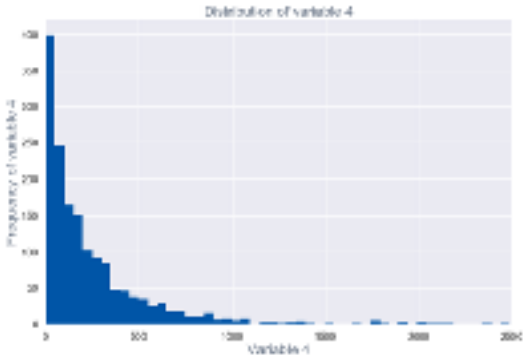

Univariate

understanding a single variable

i.e.: histogram, densityplot, barplot

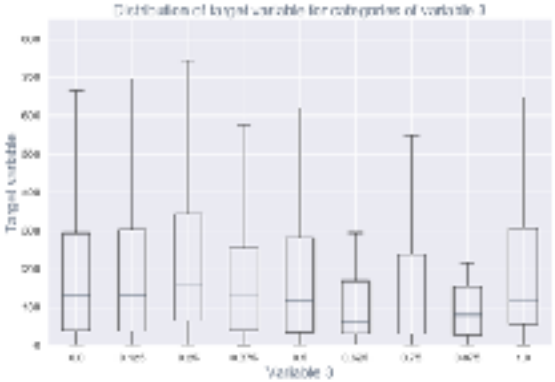

Bivariate

understanding relationship between 2 variables

i.e.: boxplot, scatterplot, grouped barplot, boxplot

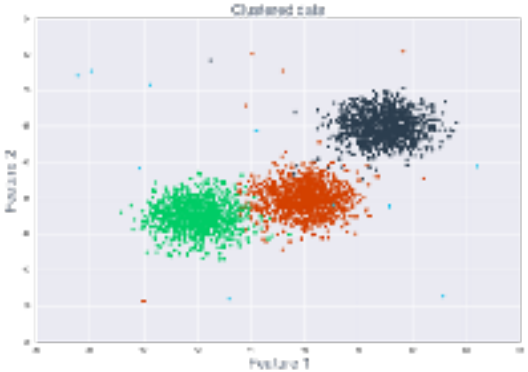

Dimensionality Reduction

projecting high-D data into a lower-D space

i.e.: PCA, ICA, Clustering

Image source: https://www.svds.com/value-exploratory-data-analysis/

```
>>> ax = sns.boxplot(x="day", y="total_bill", hue="time",
...                  data=tips, linewidth=2.5)
```

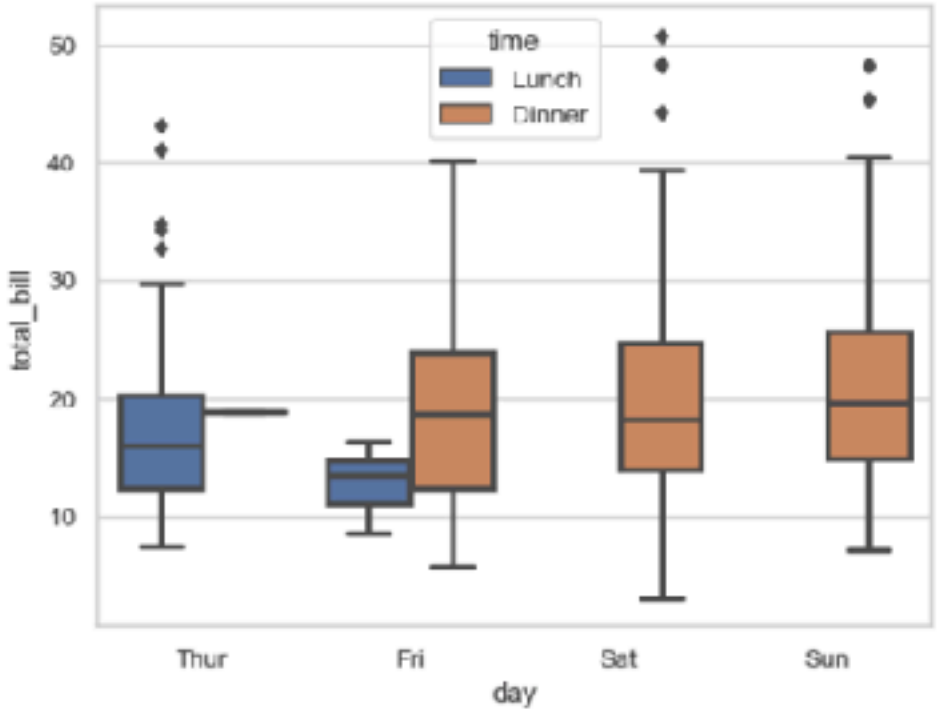

```
sns.pairplot(penguins, hue="species")
```

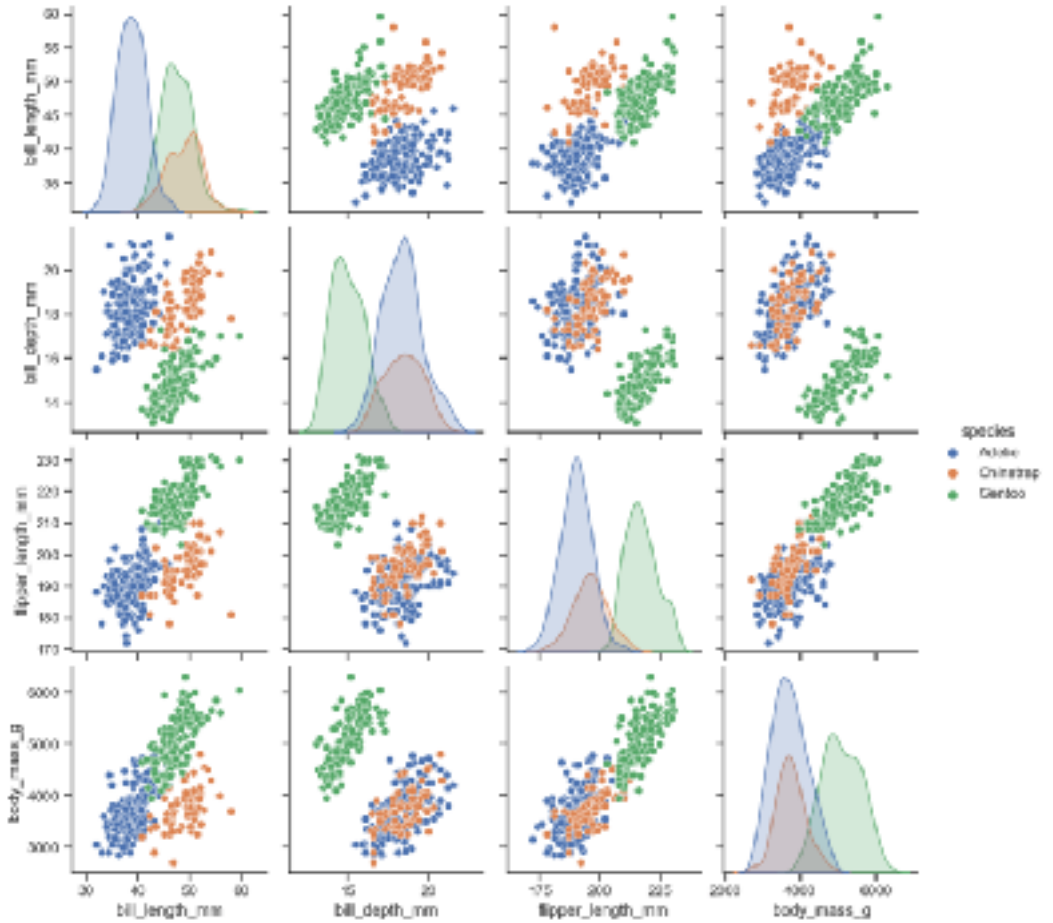

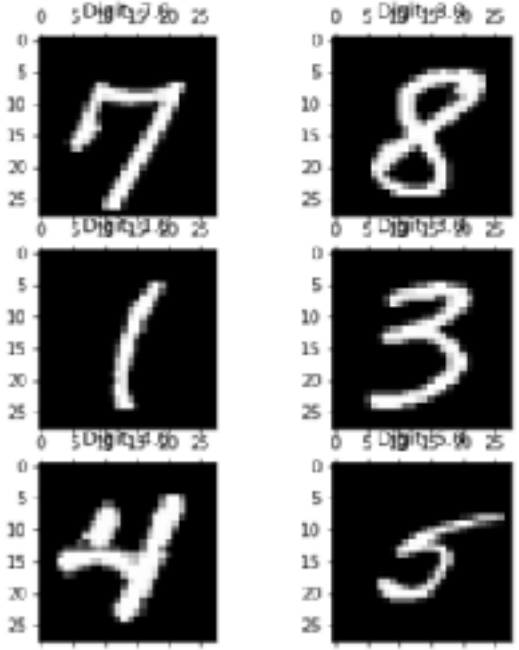

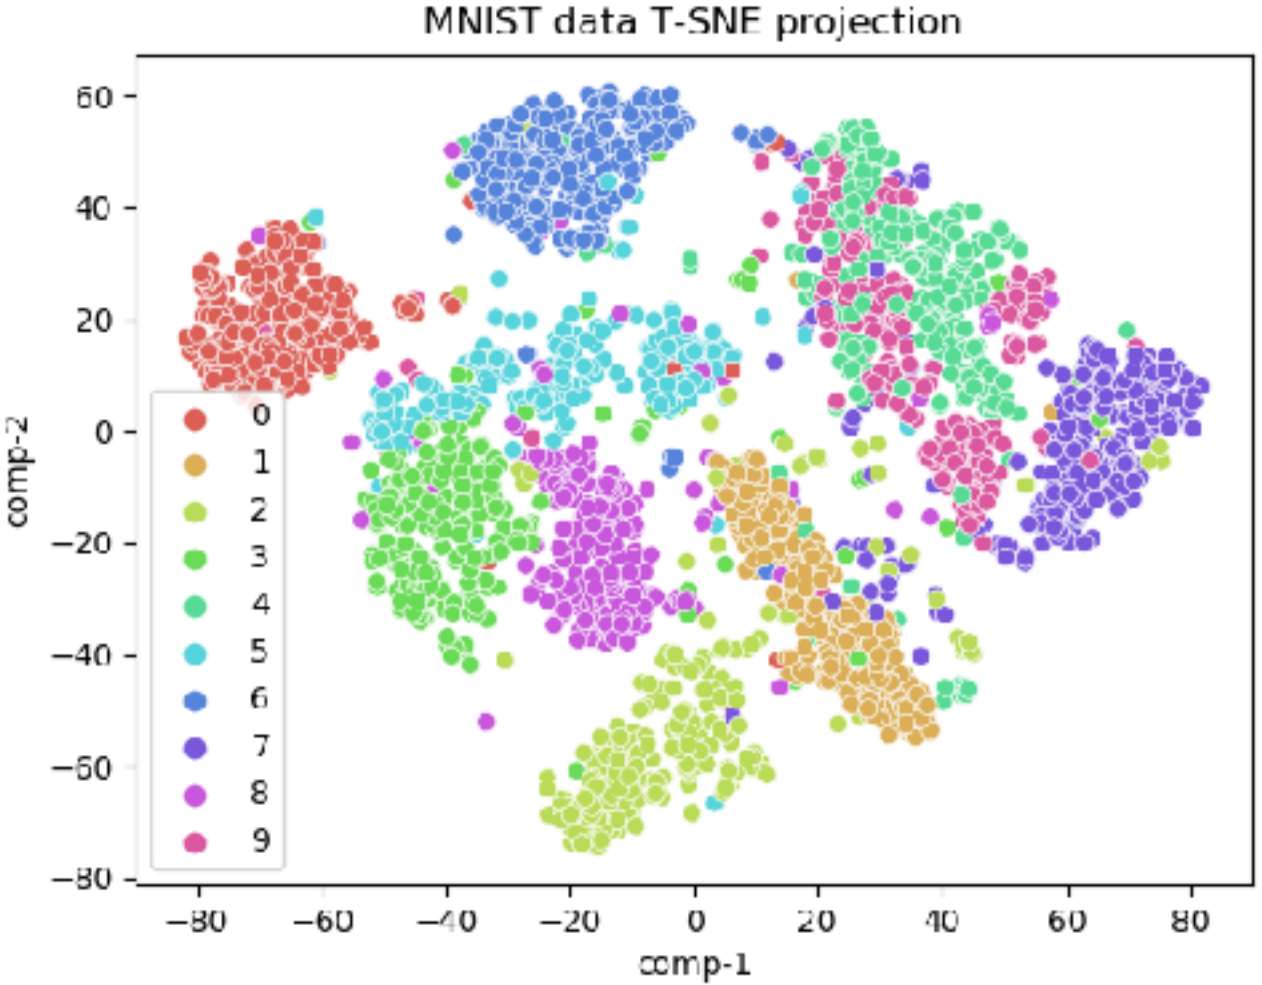
MNIST data T-SNE projection
comp-1
comp-2
0
1
2
3
4
5
6
7
8
9